

Deep Learning Architectures

Dense · Convolutional · Recurrent · Transformer

Shared objective

$$\hat{y} = F_{\theta}(x)$$

The right model depends on the structure assumed in the data.


Dense

Dense architecture

Mathematical model

$$h^{(\ell)} = \phi(W^{(\ell)}h^{(\ell-1)} + b^{(\ell)})$$
$$W^{(\ell)} \in \mathbb{R}^{m_{\ell} \times m_{\ell-1}}, b^{(\ell)} \in \mathbb{R}^{m_{\ell}}.$$

Principal layers



Inductive bias

Flat vector representation. Each output unit can combine every feature from the previous layer; the model has almost no structural prior.

Representative applications

tabular prediction regression risk scoring baseline models

Core limitation: structure must be engineered into the input.

Convolutional


Convolutional architecture

Mathematical model

$$Y_{i,j,k} = \phi\left(b_k + \sum_{u,v,c} K_{u,v,c,k} X_{i+u,j+v,c}\right)$$

A shared filter K scans local neighborhoods of the input grid.

Principal layers



Inductive bias

Locality, parameter sharing, and translation equivariance. Nearby pixels or grid cells are assumed to form reusable patterns.

Representative applications

image classification object detection segmentation

medical imaging spectrograms

Core limitation: global context needs depth, pooling, dilation, or attention.

Recurrent


Recurrent architecture

Mathematical model

$$h_t = \phi(W_x x_t + W_h h_{t-1} + b_h), \quad \hat{y}_t = g(W_y h_t + b_y)$$

LSTM/GRU variants add gates to control memory and gradients.

Principal layers



Inductive bias

Temporal order and stateful memory. The hidden state summarizes previous elements before processing the current input.

Representative applications

time series speech language modeling biomedical signals control

Core limitation: sequential state limits parallelism and gradient stability.


Transformer

Transformer architecture

Mathematical model

$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V$$
$$\text{Attn}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Principal layers



Inductive bias

Global pairwise interaction. Any token, patch, or element may attend to any other element in the input.

Representative applications

LLMs translation vision transformers multimodal AI code

Core limitation: vanilla attention scales as $O(n^2)$ with sequence length.

Similarities, Differences, and Selection Rule

Choose the architecture whose inductive bias matches the structure of the data.

Architecture	Best suited for	Main operation	Structural assumption	Main trade-off
Dense	Flat numerical or categorical feature vectors	Matrix multiplication + nonlinearity	All features may interact directly, but no spatial/temporal prior	Universal and simple, but parameter-heavy for structured data
CNN	Images, grids, spatial signals, spectrograms	Convolution + pooling	Local patterns repeat across positions	Efficient spatial feature extraction, but limited native long-range context
RNN	Ordered sequences and temporal signals	Recurrence through hidden state	Past elements influence future elements	Handles variable length, but training is sequential and less parallel
Transformer	Text, long sequences, images as patches, multimodal data	Scaled dot-product self-attention	Every element may be related to every other element	Powerful global context, but expensive for long inputs

Vector → Dense

Grid → CNN

Ordered stream → RNN

Global relations → Transformer